

A Brief Overview of Gibbs Sampling

Eric C. Rouchka
 Washington University
 Institute for Biomedical Computing
 Statistics Study Group
 May 20, 1997

Gibbs Sampler

The purpose of this overview is to present the ideas of Gibbs sampling in terms of the data, parameters, model, and procedure both in a general sense and through an application of Gibbs sampling for multiple sequence alignment.

The first requirement for the Gibbs sampler is the observable data. The observed data will be denoted Y . In the general case of the Gibbs sampler, the observed data remains constant throughout.

Gibbs sampling requires a vector of parameters of interest that are initially unknown. These parameters will be denoted by the vector Φ . Nuisance parameters, Θ , are also initially unknown. The goal of Gibbs sampling is to find estimates for the parameters of interest in order to determine how well the observable data fits the model of interest, and also whether or not data independent of the observed data fits the model described by the observed data.

Gibbs sampling requires an initial starting point for the parameters. Then, one at a time, a value for each parameter of interest is sampled given values for the other parameters and data. Once all of the parameters of interest have been sampled, the nuisance parameters are sampled given the parameters of interest and the observed data. At this point, the process is started over. The power of Gibbs sampling is that the joint distribution of the parameters will converge to the joint probability of the parameters given the observed data.

Explanation in Mathematical Terms

The Gibbs sampler requires a random starting point of parameters of interest, Φ , and nuisance parameters, Θ , with observed data Y , from which a converging distribution can be found. For the sampler, there is an initial starting point $(\Theta_1^{(0)}, \Theta_2^{(0)}, \dots, \Theta_D^{(0)}, \Phi^{(0)})$. Steps a-d are then repeatedly run.

- a) Sample $\Theta_1^{(i+1)}$ from $p(\Theta_1 | \Theta_2^{(i)}, \dots, \Theta_D^{(i)}, \Phi^{(i)}, Y)$.
- b) Sample $\Theta_2^{(i+1)}$ from $p(\Theta_2 | \Theta_1^{(i+1)}, \Theta_3^{(i)}, \dots, \Theta_D^{(i)}, \Phi^{(i)}, Y)$.
- c) Sample $\Theta_D^{(i+1)}$ from $p(\Theta_D | \Theta_1^{(i+1)}, \dots, \Theta_{(D-1)}^{(i+1)}, \Phi^{(i)}, Y)$.
- d) Sample $\Phi^{(i+1)}$ from $p(\Phi | \Theta_1^{(i+1)}, \dots, \Theta_D^{(i+1)}, Y)$.

The vectors $\Theta^{(0)}, \Theta^{(1)}, \dots, \Theta^{(t)}$ represent the realization of a Markov chain, where the transition probability from Θ' to Θ is defined as:

$$K(\Theta', \Theta) = p(\Theta_1 | \Theta_2', \dots, \Theta_D', \Phi', Y) * p(\Theta_2 | \Theta_1, \Theta_3', \dots, \Theta_D', \Phi', Y) \dots * p(\Theta_D | \Theta_1, \dots, \Theta_{(D-1)}, \Phi', Y)$$

The joint distribution of $(\Theta_1^{(i)}, \dots, \Theta_D^{(i)}, \Phi^{(i)})$ converges geometrically to $p(\Theta_1, \dots, \Theta_D, \Phi | Y)$ as $i \rightarrow \infty$.

The Gibbs sampler differs from the Metropolis algorithm because in each step only one parameter, Θ_d , is allowed to change.

Multiple Alignment Using Gibbs Sampling

One application of Gibbs sampling useful in computational molecular biology is the detection and alignment of locally conserved regions (motifs) in sequences of amino acids or nucleic acids assuming no prior information in the patterns or motifs. Gibbs sampling strategies claim to be fast and sensitive, avoiding the problem that EM algorithms fall into as far as getting trapped by local optima. As an example, a set of 29 DNA sequences have been provided. These sequences contain sequences necessary for recognition by erythroid transcription factors, most notably a six nucleotide GATA binding site.

Basic Algorithm

First the basic multiple alignment strategy is examined where a single motif is desired. The most basic implementation, known as a site sampler, assumes that there is exactly one motif element located within each sequence.

Notation

- N : number of sequences
- $S_1 \dots S_N$: set of sequences
- W : width of motif to be found in the sequences
- J : the number of residues in the alphabet. $J = 4$ for nucleic acid sequences and 20 for amino acid sequences.
- $c_{i,j,k}$: Observed counts of residue j in position i of motif k . j ranges from $1 \dots J$. i ranges from $0..W$ where $c_{0,j}$ contains the counts of residue j in the background. If it is assumed that only a single motif is searched for, the k term can drop out.
- $q_{i,j}$: frequency of residue j occurring in position i of the motif. i ranges from $0..W$ as above. Note that in the literature, $q_{0,j}$ (the vector of background residue frequencies) is sometimes denoted as p_j . This is the parameter of interest, Φ .
- a_k : vector of starting positions of the motifs within the sequences. k ranges from $1..N$. This is the nuisance parameter, Θ .
- b_j : pseudocounts for each residue – needed according to Bayesian statistical rules to eliminate problems with zero counts.
- B : The total number of pseudocounts. $B = \sum_j b_j$.

Initialization

Once the sequences are known, the counts for each residue can be calculated. Initially, $c_{0,j}$ will contain the total counts of residue j within all of the sequences and $c_{i,j}$ is initialized to 0 for all other values of i . This is a summary of observed data. The site sampler is then initialized by randomly selecting a position for the motif within each sequence and recording these positions in a_k . The counts are updated according to this initial alignment. After the observed counts are set, $q_{i,j}$ can be calculated.

$$q_{i,j} = \frac{c_{i,j} + b_j}{N - 1 + B}$$

Equation 1: Motif Residue Frequencies

$$q_{0,j} = \frac{c_{0,j} + b_j}{\sum_{k=1}^j c_{0,k} + B}$$

Equation 2 : Background Residue Frequencies

Predictive Update Step

The first step, known as the predictive update step, selects one of the sequences and places the motif within that sequence in the background and updates the residue counts. One of the N sequences, z , is chosen. The motif in sequence z is taken from the model and placed in the background. The observed counts $c_{i,j}$ are updated as are the frequencies $q_{i,j}$. The selection of z can be random or in a specified order.

Sampling Step

In the sampling step, a new motif position for the selected sequence is determined by sampling according to a weight distribution. All of the possible segments of width W within sequence z are considered. For each of these segments x , a weight A_x is

calculated according to the ratio $A_x = \frac{Q_x}{P_x}$ where $Q_x = \prod_{i=1}^W q_{i,r_i}$ is the model residue

frequency according to equation 1 if segment x is in the motif model, and

$P_x = \prod_{i=1}^W q_{0,r_i}$ is the background residue frequency according to equation 2. r_i refers to

the residue located at position i of segment x . Once A_x is calculated for every possible x , a new position a_z is chosen by randomly sampling over the set of weights A_x . Thus, possible starting positions with higher weights will be more likely to be chosen as the new motif position than those position with lower weights. Since this is a stochastic process, the starting position with the highest weight is not guaranteed to be chosen.

Once the iterative predictive update and sampling steps have been performed for all of the sequences, a probable alignment is present. For this alignment, a maximum a posteriori (MAP) estimate can be calculated using equation 3:

$$F = \sum_{i=1}^W \sum_{j=1}^J c_{i,j} \log \frac{q_{i,j}}{q_{0,j}}$$

Equation 1: Alignment conditional log-likelihood

The goal is to maximize F . This is accomplished in the following manner:

```
globalMaxAlignmentProb = 0
For Iteration = 1 to 10:
    Initialize Random alignment
    localMaxAlignmentProb = 0;
    while (not in local maximum and innerloop < MAXLOOP) do
        for each sequence do{
            Predictive Update
            Sample
        }
        calculate AlignmentProb
        if (AlignmentProb > localMaxAlignmentProb)
        {
            localMaxAlignmentProb = AlignmentProb;
            not in local maximum = true;
        }
        innerloop++;
    }
    if (localMaxAlignmentProb == globalMaxAlignmentProb)
        exit → max found twice
    else if (localMaxAlignmentProb > globalMaxAlignmentProb)
        globalMaxAlignmentProb = localMaxAlignmentProb
}
```

Explanation

The idea is that the more accurate the predictive update step is, the more accurate the sampling step will be since the background will be more distinguished from the motif description. Given random positions a_k in the sampling step, the pattern description $q_{i,j}$ will not favor any particular segment. Once some correct a_k have been selected by chance, the $q_{i,j}$ begins to favor a particular motif.

Details

There are a couple of problems that need to be addressed. First, it is possible that the correct pattern has not been chosen, but rather a shift of it has. This can be taken care of by shifting the alignment to the left and right by a specified number of columns and sampling from the values of F .

Another problem is that the pattern width W must also be specified. In order to decide what the width should be, the incomplete-data log-probability ratio as shown in Equation 4 can be implemented.

$$G = F - \sum_{i=1}^N (\log L'_i + \sum_{j=1}^{L'_i} Y_{i,j} \log Y_{i,j})$$

Equation 2 : Incomplete-data log-probability ratio

Where L'_i is the number of the possible positions for the pattern within sequence i and $Y_{i,j}$ is the normalized weight of position j . Dividing G by the number of free parameters needed to specify the pattern ($19 * W$ for protein sequences, $3 * W$ for nucleotide sequences) results in an information per parameter quantity. It is then desired to maximize the information per parameter to determine the value of W .

Improving the Algorithm

The method of determining motifs as described above requires multiple runs on the same data set with varying widths to find the correct pattern size. The *Protein Science* paper discusses a method to determine the width of the motif in a single run of the program, while at the same time determining gaps within the motif. The Gibbs sampler described thus far also requires the existence of exactly one motif in each sequence. Another improvement made is to allow multiple motifs within sequences, and allow the possibility that a sequence does not have any motifs. The improvements made within the *Protein Science* paper describe a technique known as a motif sampler.

Allowing a Variable Number of Motif Sites

Assume that there are m different motif patterns that we are searching for in the sequences. Let n_k represent the number of sites matching motif k in the sequence. Initially, it is not known how many motif sites there are. To overcome this, we make a prior expectation e_k for each n_k . The new algorithm allows the prior expectations to become posterior expectations as it learns the number of sites for each motif. For the

initialization step, e_k random starting points are selected for motif pattern k instead of selecting one starting point randomly within each of the N sequences. Now we can go through all of the possible motif starting locations in each sequence and decide if it is a motif starting site by using equation 5.

$$\frac{p_j}{1 - p_j} A_x$$

Equation 3: Current motif site probability

Where p_j is the posterior probability that any site belongs to the model (see the appendix of the *Protein Science* paper for the prior and posterior calculations of P_j), and A_x is the same as in the site sampler.

Width Optimization by Column Sampling

In order to help introduce gaps and include only the most informative positions of the motif, column sampling is introduced where only C columns out of a specified number of contiguous columns $w_{\max} \geq C$ are used for the residue frequency model. This is accomplished in a two step process. First, turn off one column either randomly or by selecting it proportional to how little information it provides. Then sample one of the columns that are turned off proportional to how information rich it is and turn it on. The column move operations need to be weighted in order to assure that there is not a bias to longer motif widths. A discussion is provided in the appendix of the *Protein Science* paper.

Properties of Gibbs Sampling

- Requires relatively large sets (15 or more sequences) for weakly conserved patterns to reach statistical significance.
- It is a heuristic and not an exhaustive search, so you are not guaranteed to reach an optimal value, but you will not get stuck in local maximums the way EM algorithms do.
- Have to have some idea of how wide the motif is and how many motifs there are for the algorithm to work best.
- Gibbs sampling allows you to view suboptimal results.
- Fast and sensitive – generally finds an optimized local alignment model for N -sequences in N -linear time.

Site Sampler Example

The site sampler is tested using a set of erythroid sequences. The set is tested for the presence of a GATA box, which should have a sequence (T/A)GATA(A/G), which in the reverse complement is (C/T)TATC(A/T). Since the width of the GATA box is shown, it is known that for this example $W = 6$. The process of determining the best alignment using the site sampler is described.

Initialization

The first step in the site sampler is to randomly assign an alignment to the set of sequences. Figure 1 indicates one such random alignment.

```

TCAGAACCAGTTATAAAATTTATCATTTCCTTCTCCACTCCT
CCCACGCAGCCGCCCTCCTCCCCGGTCACTGACTGGTCCTG
TCGACCCTCTGAACCTATCAGGGACCACAGTCAGCCAGGCAAG
AAACACTTGAGGGAGCAGATAACTGGGCCAACCATGACTC
GGGTGAATGGTACTGCTGATTACAACCTCTGGTGCTGC
AGCCTAGAGTGATGACTCCTATCTGGGTCCCCAGCAGGA
GCCTCAGGATCCAGCACACATTATCACAAACTTAGTGTCCA
CATTATCACAAACTTAGTGTCCATCCATCACTGCTGACCCT
TCGGAACAAGGCAAAGGCTATAAAAAAATTAAGCAGC
GCCCTTCCCCACACTATCTCAATGCAAATATCTGTCTGAAACGGTTCC
CATGCCCTCAAGTGTGCAGATTGGTCACAGCATTTCAAGG
GATTGGTCACAGCATTTCAAGGGAGAGACCTCATTGTAAG
TCCCCAACTCCCAACTGACCTTATCTGTGGGGGAGGCTTTTGA
CCTTATCTGTGGGGGAGGCTTTTGAAGTAATTAGGTTTAGC
ATTATTTTCCTTATCAGAAGCAGAGAGACAAGCCATTTCTCTTTCCCTCCCGGT
AGGCTATAAAAAAAATTAAGCAGCAGTATCCTCTTGGGGGCCCTTC
CCAGCACACACACTTATCCAGTGGTAAATACACATCAT
TCAAATAGGTACGGATAAGTAGATATTGAAGTAAGGAT
ACTTGGGGTTCCAGTTTGATAAGAAAAGACTTCCTGTGGA
TGGCCGCAGGAAGGTGGGCCTGGAAGATAACAGCTAGTAGGCTAAGGCCAG
CAACCACAACCTCTGTATCCGGTAGTGGCAGATGGAAA
CTGTATCCGGTAGTGGCAGATGGAAAGAGAAACGGTTAGAA
GAAAAAAAATAAATGAAGTCTGCCTATCTCCGGGCCAGAGCCCCCT
TGCCTTGTCTGTTGTAGATAATGAATCTATCCTCCAGTGACT
GGCCAGGCTGATGGGCCTTATCTCTTTACCCACCTGGCTGT
CAACAGCAGGTCCTACTATCGCCTCCCTCTAGTCTCTG
CCAACCGTTAATGCTAGAGTTATCACTTTCTGTTATCAAGTGGCTTCAGCTATGCA
GGGAGGGTGGGGCCCCTATCTCTCCTAGACTCTGTG
CTTTGTCAACTGGATCTGATAAGAAACACCACCCCTGC

```

Figure 1: Initial Motif Sites for Site Sampler

The number of A's in all of the sequences combined is 327, the number of C's is 317, the number of G's is 272, and the number of T's is 304. If we assume that the pseudocounts

needed for Bayesian statistics are equal to 10% of the observed nucleotide frequencies, we know the following information before any of the initial motif sites are set:

$$c_{0,1} = 327; \quad c_{0,2} = 317; \quad c_{0,3} = 272; \quad c_{0,4} = 304; \quad \sum_{i=1}^4 c_{0,i} = 1220$$

$$b_1 = 32.7; \quad b_2 = 31.7; \quad b_3 = 27.2; \quad b_4 = 30.4; \quad B = 122$$

If we assume we have the initial random alignment as described in figure 1, we can recalculate the counts and calculate the residue frequencies. Table 1 gives the results of these calculations.

| Nucleotide | Motif Position (0 = Background) | | | | | | |
|------------|---------------------------------|---|----|----|---|----|---|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| A | 279 | 6 | 12 | 6 | 6 | 11 | 7 |
| C | 280 | 8 | 3 | 5 | 7 | 7 | 7 |
| G | 225 | 9 | 8 | 10 | 7 | 5 | 8 |
| T | 262 | 6 | 6 | 8 | 9 | 6 | 7 |

Table 1: Calculation of observed counts for initial alignment of figure 1

| Nucleotide | Motif Position (0 = Background) | | | | | | |
|------------|---------------------------------|------|------|------|------|------|------|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| A | .267 | .256 | .296 | .256 | .256 | .289 | .263 |
| C | .267 | .263 | .230 | .243 | .256 | .256 | .256 |
| G | .216 | .240 | .233 | .246 | .226 | .213 | .233 |
| T | .250 | .241 | .241 | .254 | .261 | .241 | .248 |

Table 2: Calculation of residue frequencies for initial alignment of figure 1

Predictive Update Step

Now that we have the initial random alignment for the site sampler, we begin the predictive update step by choosing one of the sequences to update. For simplicity, let's just choose the first sequence. In the predictive update stage, the motif for the selected

sequence is placed in the background and the counts and frequencies are updated. Since the motif in the first sequence is ATTTAT, tables 1 and 2 can be recalculated.

| Nucleotide | Motif Position (0 = Background) | | | | | | |
|------------|---------------------------------|---|----|----|---|----|---|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| A | 281 | 5 | 12 | 6 | 6 | 10 | 7 |
| C | 280 | 8 | 3 | 5 | 7 | 7 | 7 |
| G | 225 | 9 | 8 | 10 | 7 | 5 | 8 |
| T | 266 | 6 | 5 | 5 | 5 | 6 | 6 |

Table 3: Recalculated observed counts

| Nucleotide | Motif Position (0 = Background) | | | | | | |
|------------|---------------------------------|------|------|------|------|------|------|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| A | .267 | .251 | .298 | .258 | .258 | .285 | .265 |
| C | .265 | .264 | .231 | .245 | .258 | .258 | .258 |
| G | .215 | .241 | .235 | .248 | .228 | .215 | .235 |
| T | .252 | .243 | .236 | .236 | .236 | .243 | .243 |

Table 4: Recalculated residue frequencies

Sampling Step

Once the counts and frequencies have been updated in the predictive update step, we can begin the sampling step. In this step, we look at all of the possible motif starting positions within the sequence selected in the predictive update step. Since the length of the sequence is 41 and the width of the motif is 6, there are $41 - 6 + 1 = 36$ possible starting sites. We calculate the probability of each of these sites being in the model and then sample from their weights. Table 5 shows the weights for sequence 1 given table 3 and 4.

Using the information in table 5, one of the segments will be sampled in according to the normalized value of A_x . The predictive update and sampling steps are repeated for each of the sequences. Once each of the sequences have been sampled, an alignment is present and the alignment probability is tested. This procedure is repeated until a plateau is reached. Then another initial random alignment is tested and the process begins again. For the example used thus far, the final alignment is as shown in figure 2. This alignment yields the counts and frequencies described in tables 6 and 7.

| Sequence | X | A_x | Normalized A_x |
|-----------------|----------|----------------------|---------------------------------|
| TCAGAA | 1 | 0.906 | .030 |
| CAGAAC | 2 | 1.283 | .042 |
| AGAACC | 3 | 0.905 | .030 |
| GAACCA | 4 | 1.131 | .037 |
| AACCAG | 5 | 1.093 | .036 |
| ACCAGT | 6 | 0.700 | .023 |
| CCAGTT | 7 | 0.822 | .027 |
| CAGTTA | 8 | 1.142 | .037 |
| AGTTAT | 9 | 0.919 | .030 |
| GTTATA | 10 | 0.902 | .029 |
| TTATAA | 11 | 0.856 | .028 |
| TATAAA | 12 | 1.021 | .033 |
| ATAAAT | 13 | 0.839 | .027 |
| TAAATT | 14 | 0.923 | .030 |
| AAATTT | 15 | 0.875 | .029 |
| AATTTA | 16 | 0.873 | .028 |
| ATTTAT | 17 | 0.787 | .026 |
| TTTATC | 18 | 0.757 | .025 |
| TTATCA | 19 | 0.781 | .025 |
| TATCAT | 20 | 0.997 | .033 |
| ATCATT | 21 | 0.723 | .024 |
| TCATTT | 22 | 0.698 | .023 |
| CATTTC | 23 | 0.907 | .030 |
| ATTTCC | 24 | 0.725 | .024 |
| TTTCCT | 25 | 0.762 | .025 |
| TTCCTT | 26 | 0.743 | .024 |
| TCCTTC | 27 | 0.674 | .022 |
| CCTTCT | 28 | 0.709 | .023 |
| CTTCTC | 29 | 0.791 | .026 |
| TTCTCC | 30 | 0.731 | .024 |
| TCTCCA | 31 | 0.732 | .024 |
| CTCCAC | 32 | 0.864 | .028 |
| TCCACT | 33 | 0.696 | .023 |
| CCACTC | 34 | 0.761 | .025 |
| CACTCC | 35 | 0.904 | .030 |
| ACTCCT | 36 | 0.695 | .023 |

Table 5 :Weights for segments within sequence 1

TCAGAACCAGTTATAAAT**TTATCA**TTTCCTTCTCCACTCCT
 CCCACGCAGCCGCCCTCCTCCC**CGGTCA**CTGACTGGTCCTG
 TCGACCCCTCTGGAAC**CTATCA**GGGACCACAGTCAGCCAGGCAAG
 AAAACACTTGAGGGAGC**AGATAA**CTGGGCCAACCATGACTC
 GGGTGAA**TGGTAC**TGCTGATTACAACCTCTGGTGCTGC
 AGCCTAGAGTGATGACTC**CTATCT**GGGTCCCCAGCAGGA
 GCCTCAGGATCCAGCACACA**TTATCA**CAAACCTTAGTGTCCA
 CA**TTATCA**CAAACCTTAGTGTCCATCCATCACTGCTGACCCCT
 TCGGAACAAGGCAAAGG**CTATAA**AAAAAATTAAGCAGC
 GCCCCCTTCCCCACA**CTATCT**CAATGCAAATATCTGTCTGAAACGGTTCC
 CATGCCCTCAAGTGTGCAGAT**TGGTCA**CAGCATTTCAAGG
 GATTGGTACACAGCATTTCAAGGGAGAGACCTCA**TTGTAAG**
 TCCCCAACTCCCAACTGACC**TTATCT**GTGGGGGAGGCTTTTGA
 CCT**TTATCT**GTGGGGGAGGCTTTTGAAAAGTAATTAGGTTTAGC
 ATTATTTTCC**TTATCA**GAAAGCAGAGAGACAAGCCATTTCTCTTTCCTCCCGGT
 AGG**CTATAA**AAAAAATTAAGCAGCAGTATCCTCTTGGGGGCCCCCTTC
 CCAGCACACACAC**TTATCC**AGTGGTAAATACACATCAT
 TCAAATAGGTACGGATAAGT**AGATAT**TGAAGTAAGGAT
 ACTTGGGGTTCCAGTT**TGATAA**GAAAAGACTTCCTGTGGA
 TGGCCGCAGGAAGGTGGGCCTGGA**AGATAA**CAGCTAGTAGGCTAAGGCCAG
 CAACCACAACCT**CTGTAT**CCGGTAGTGGCAGATGGAAA
CTGTATCCGGTAGTGGCAGATGGAAAGAGAAACGGTTAGAA
 GAAAAAAAATAAATGAAGTCTGC**CTATCT**CCGGGCCAGAGCCCCT
 TGCC**TTGTCT**GTTGTAGATAATGAATCTATCCTCCAGTGA
 GGCCAGGCTGATGGGCC**TTATCT**CTTTACCCACCTGGCTGT
 CAACAGCAGGTCCTACTATCGCCTCCCTC**TAGTCT**CTG
 CCAACCGTTAATGCTAGAGTTATCACTTTCTG**TTATCA**AGTGGCTTCAGCTATGCA
 GGGAGGGTGGGGCCCC**CTATCT**CTCCTAGACTCTGTG
 CTTTGTCACTGGATC**TGATAA**GAAACACCACCCCTGC

Figure 2: Final Alignment for Site Sampler

| Nucleotide | Motif Position (0 = Background) | | | | | | |
|------------|---------------------------------|----|----|----|----|----|----|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| A | 276 | 3 | 1 | 21 | 0 | 11 | 15 |
| C | 287 | 10 | 0 | 0 | 0 | 18 | 2 |
| G | 256 | 0 | 8 | 8 | 0 | 0 | 0 |
| T | 227 | 16 | 20 | 0 | 29 | 0 | 12 |

Table 6: Final site sampler residue counts

| Nucleotide | Motif Position (0 = Background) | | | | | | |
|-------------------|----------------------------------------|----------|----------|----------|----------|----------|----------|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 |
| A | .264 | .236 | .223 | .356 | .217 | .289 | .315 |
| C | .273 | .276 | .210 | .210 | .210 | .329 | .223 |
| G | .242 | .180 | .233 | .233 | .180 | .180 | .180 |
| T | .220 | .307 | .334 | .201 | .393 | .201 | .281 |

Table 7: Final site sampler residue frequencies

If we were to go back to the predictive update/sampling stages with these results, we would sample from the results gathered in table 8.

| Sequence | X | A_x | Normalized A_x |
|----------|----|-------|------------------|
| TCAGAA | 1 | 1.369 | .023 |
| CAGAAC | 2 | 0.631 | .011 |
| AGAACC | 3 | 0.941 | .016 |
| GAACCA | 4 | 0.922 | .016 |
| AACCAG | 5 | 0.380 | .006 |
| ACCAGT | 6 | 0.427 | .007 |
| CCAGTT | 7 | 0.921 | .016 |
| CAGTTA | 8 | 1.591 | .027 |
| AGTTAT | 9 | 1.968 | .034 |
| GTTATA | 10 | 0.903 | .015 |
| TTATAA | 11 | 6.114 | .104 |
| TATAAA | 12 | 1.129 | .019 |
| ATAAAT | 13 | 2.048 | .035 |
| TAAATT | 14 | 1.475 | .025 |
| AAATTT | 15 | 2.071 | .035 |
| AATTTA | 16 | 1.316 | .022 |
| ATTTAT | 17 | 3.004 | .051 |
| TTTATC | 18 | 1.150 | .020 |
| TTATCA | 19 | 6.622 | .113 |
| TATCAT | 20 | 1.152 | .020 |
| ATCATT | 21 | 0.994 | .017 |
| TCATTT | 22 | 2.873 | .049 |
| CATTTC | 23 | 1.048 | .018 |
| ATTTCC | 24 | 2.106 | .036 |
| TTTCCT | 25 | 2.183 | .037 |
| TTCCTT | 26 | 1.421 | .024 |
| TCCTTC | 27 | 1.087 | .019 |
| CCTTCT | 28 | 1.931 | .033 |
| CTTCTC | 29 | 0.810 | .014 |
| TTCTCC | 30 | 2.715 | .046 |
| TCTCCA | 31 | 1.053 | .018 |
| CTCCAC | 32 | 0.829 | .014 |
| TCCACT | 33 | 1.029 | .018 |
| CCACTC | 34 | 0.615 | .010 |
| CACTCC | 35 | 1.161 | .020 |
| ACTCCT | 36 | 0.751 | .013 |

Table 8: Weights for sampling step after a near maximal alignment has been found

Comparison of site sampler with the motif sampler

The site sampler and motif sampler follow the same basic Gibbs techniques. The difference is that the motif sampler will allow for the detection of zero or more motif locations in each sequence, whereas the site sampler detects exactly one. Thus, for the initialization step with the motif sampler, a random alignment is made according to a guess estimate as to how many motif sites exist in total. An example of an initial alignment is given in figure 3.

```

TCAGAACCAGTTATAAAATTTATCATTTTCCTTCTCCACTCCT
CCCACGCAGCCGCCCTCCTCCCCGGTCACTGACTGGTCCTG
TCGACCCTCTGAACCTATCAGGGACCACAGTCAGCCAGGCAAG
AAACACCTTGAGGGAGCAGATAACTGGGCCAACCATGACTC
GGGTGAATGGTACTGCTGATTACAACCTCTGGTGCTGC
AGCCTAGAGTGATGACTCCTATCTGGGTCCCCAGCAGGA
GCCTCAGGATCCAGCACACATTATCACAACTTAGTGTCCA
CATTATCACAACTTAGTGTCCATCCATCACTGCTGACCCT
TCGGAACAAGGCAAAGGCTATAAAAAAATTAAGCAGC
GCCCCCTTCCCCACACTATCTCAATGCAAATATCTGTCTGAAACGGTTCC
CATGCCCTCAAGTGTGCAGATTGGTCACAGCATTTCAGG
GATTGGTCACAGCATTTCAGGGAGAGACCTCATTGTAAG
TCCCCAACTCCCAACTGACCTTATCTGTGGGGGAGGCTTTTGA
CCTTATCTGTGGGGGAGGCTTTTGAAGTAATTAGGTTTAGC
ATTATTTTCCTTATCAGAAGCAGAGAGACAAGCCATTTCTCTTTTCCTCCCGGT
AGGCTATAAAAAAATTAAGCAGCAGTATCCTCTTGGGGGCCCTTC
CCAGCACACACACTTATCCAGTGGTAAATACACATCAT
TCAAATAGGTACGGATAAGTAGATATTGAAGTAAGGAT
ACTTGGGGTTCCAGTTTGATAAGAAAAGACTTCCTGTGGA
TGGCCGCAGGAAGGTGGGCCTGGAAGATAACAGCTAGTAGGCTAAGGCCAG
CAACCACAACCTCTGTATCCGGTAGTGGCAGATGGAAA
CTGTATCCGGTAGTGGCAGATGGAAAGAGAAACGGTTAGAA
GAAAAAAATAAATGAAGTCTGCCTATCTCCGGGCCAGAGCCCCT
TGCCTTGTCTGTTGTAGATAATGAATCTATCCTCCAGTGA
GGCCAGGCTGATGGGCCTTATCTCTTTACCCACCTGGCTGT
CAACAGCAGGTCCTACTATCGCCTCCCTCTAGTCTCTG
CCAACCGTTAATGCTAGAGTTATCACTTTCTGTTATCAAGTGGCTTTACAGCTATGCA
GGGAGGGTGGGGCCCCCTATCTCTCCTAGACTCTGTG
CTTTGTCACTGGATCTGATAAGAAACACCACCCCTGC

```

Figure 3: Initial Alignment for Motif Sampler ($e = 30$)

Note that the `guestimate` does not need to be the exact number of motif positions to be found. This is just a starting number that will evolve within the motif sampler. Using the same data that is used with the site sampler, the maximal alignment using the motif sampler is given in figure 4. With the motif sampler, any given location is sampled into the model based on the ratio of the site being in the model to it being in the background.

```

TCAGAACCAGTTATAAATTTATCATTTCCTTCTCCACTCCT
CCCACGCAGCCGCCCTCCTCCCCGGTCACTGACTGGTCCTG
TCGACCCTCTGGAACCTATCAGGGACCACAGTCAGCCAGGCAAG
AAAACACTTGAGGGAGCAGATAACTGGGCCAACCATGACTC
GGGTGAATGGTACTGCTGATTACAACCTCTGGTGCTGC
AGCCTAGAGTGATGACTCCTATCTGGGTCCCCAGCAGGA
GCCTCAGGATCCAGCACACATTATCACAAACCTTAGTGTCCA
CATTATCACAAACCTTAGTGTCCATCCATCACTGCTGACCCT
TCGGAACAAGGCAAAGGCTATAAAAAAATTAAGCAGC
GCCCCTTCCCCACACTATCTCAATGCAAATATCTGTCTGAAACGGTTCC
CATGCCCTCAAGTGTGCAGATTGGTCAAGCATTTCAGG
GATTGGTCAAGCATTTCAGGGAGAGACCTCATTTGTAAG
TCCCCAACTCCCAACTGACCTTATCTGTGGGGGAGGCTTTTGA
CCTTATCTGTGGGGGAGGCTTTTGAAAAGTAATTAGGTTTAGC
ATTATTTTCCTTATCAGAAAGCAGAGAGACAAGCCATTTCTCTTTCTCCCGGT
AGGCTATAAAAAAATTAAGCAGCAGTATCCTCTTGGGGGCCCCCTTC
CCAGCACACACACTTATCCAGTGGTAAATACACATCAT
TCAAATAGGTACGGATAAGTAGATATTGAAGTAAGGAT
ACTTGGGGTTCCAGTTTGATAAGAAAAGACTTCCTGTGGA
TGGCCGCAGGAAGGTGGGCCTGGAAGATAACAGCTAGTAGGCTAAGGCCAG
CAACCACAACCTCTGTATCCGGTAGTGGCAGATGGAAA
CTGTATCCGGTAGTGGCAGATGGAAAGAGAAACGGTTAGAA
GAAAAAAAATAAATGAAGTCTGCCTATCTCCGGGCCAGAGCCCCCT
TGCCTTGTCTGTTGTAGATAATGAATCTATCCTCCAGTGA
GGCCAGGCTGATGGGCCTTATCTCTTTACCCACCTGGCTGT
CAACAGCAGGTCCTACTATCGCCTCCCTCTAGTCTCTG
CCAACCGTTAATGCTAGAGTTATCACTTTCTGTTATCAAGTGGCTTCAGCTATGCA
GGGAGGGTGGGGCCCCTATCTCTCCTAGACTCTGTG
CTTTGTCACTGGATCTGATAAGAAACACCACCCCTGC

```

Figure 2: Final alignment for motif sampler

References

Lawrence CE, Altschul SF, Boguski MS, Liu JS, Neuwald AF, Wootton JC. 1993. Detecting subtle sequence signals: A Gibbs sampling strategy for multiple alignment. *Science* **262**:208-214.

This paper describes a Gibbs sampling strategy where there is assumed to be a single occurrence of the motif within each sequence. No gaps are allowed within the alignment. The implementation is known in the literature as a site sampler.

Neuwald AF, Liu JS, Lawrence CE. 1995. Gibbs motif sampling: detection of outer membrane repeats. *Protein Science* **4**:1618-1632.

This paper describes a Gibbs sampling strategy where the number of motifs is not known. This is the motif sampler. Examples are presented in the location of the immunoglobulin fold and hth motifs.

Liu JS, Neuwald AF, Lawrence CE. 1995. Bayesian models for multiple local sequence alignment and Gibbs sampling strategies. *Journal of the American Statistical Association* **90**, 432:1156-1171.

This paper contains more of the derivations for implementing a Bayesian model in the Gibbs sampler. I have not covered any of the details of this paper here, but if further research into the derivations of the various formulas sounds interesting, this should be a good place to start.

Tanner, MA. 1993. Tools for Statistical Inference. Springer-Verlag.